



Seattle Airbnb Open Data



Analysis by:

- Stergios Sozos 8170129
- Aliko Ntouzgou 8170091

A sneak peek into the Airbnb activity in Seattle, WA, USA for the year 2016

Data Info



LISTINGS

Full descriptions and average score
From 2009 to 2015



REVIEWS

Reviewer ID and detailed
comments
From 2009 to 2015



CALENDAR

Listing's price and availability each
date
From 2016 to 2017

Data metrics



Rows per file



Listings: 3.818



Reviews: 84.849



Calendar: 1.393.570



Columns per file



Listings: 92



Reviews: 6



Calendar: 4

Main categories - Dimensions



LOCATION



PRICE



OPTIONS



FACILITIES



BED

Data Cleaning

Selected usable
columns

Fixed NaNs and
empty cells

Turned float into int

Renamed columns,
for better
understanding and
easier use

Turned “date”
column into
datetime type

Removed Dollar
sign from Prices



Star Scheme

	Column Name	Data Type	Allow Nulls
	neighbourhood_group_cleansed	nvarchar(50)	<input type="checkbox"/>
	latitude	float	<input type="checkbox"/>
	longitude	float	<input type="checkbox"/>
▼	location_id	int	<input type="checkbox"/>
			<input type="checkbox"/>

Foreign Key Relationships

Selected Relationship:

- FK_Listing_BedDim
- FK_Listing_DateDim
- FK_Listing_FacilitiesDim
- FK_Listing_HostDim
- FK_Listing_LocationDim
- FK_Listing_OptionsDim
- FK_Listing_PriceDim

Editing properties for existing relationship.

▼ (General)

Check Existing Data On CI Yes

► Tables And Columns Specification

▼ Identity

(Name) FK_Listing_BedDim

Description

▼ Table Designer

Enforce For Replication Yes

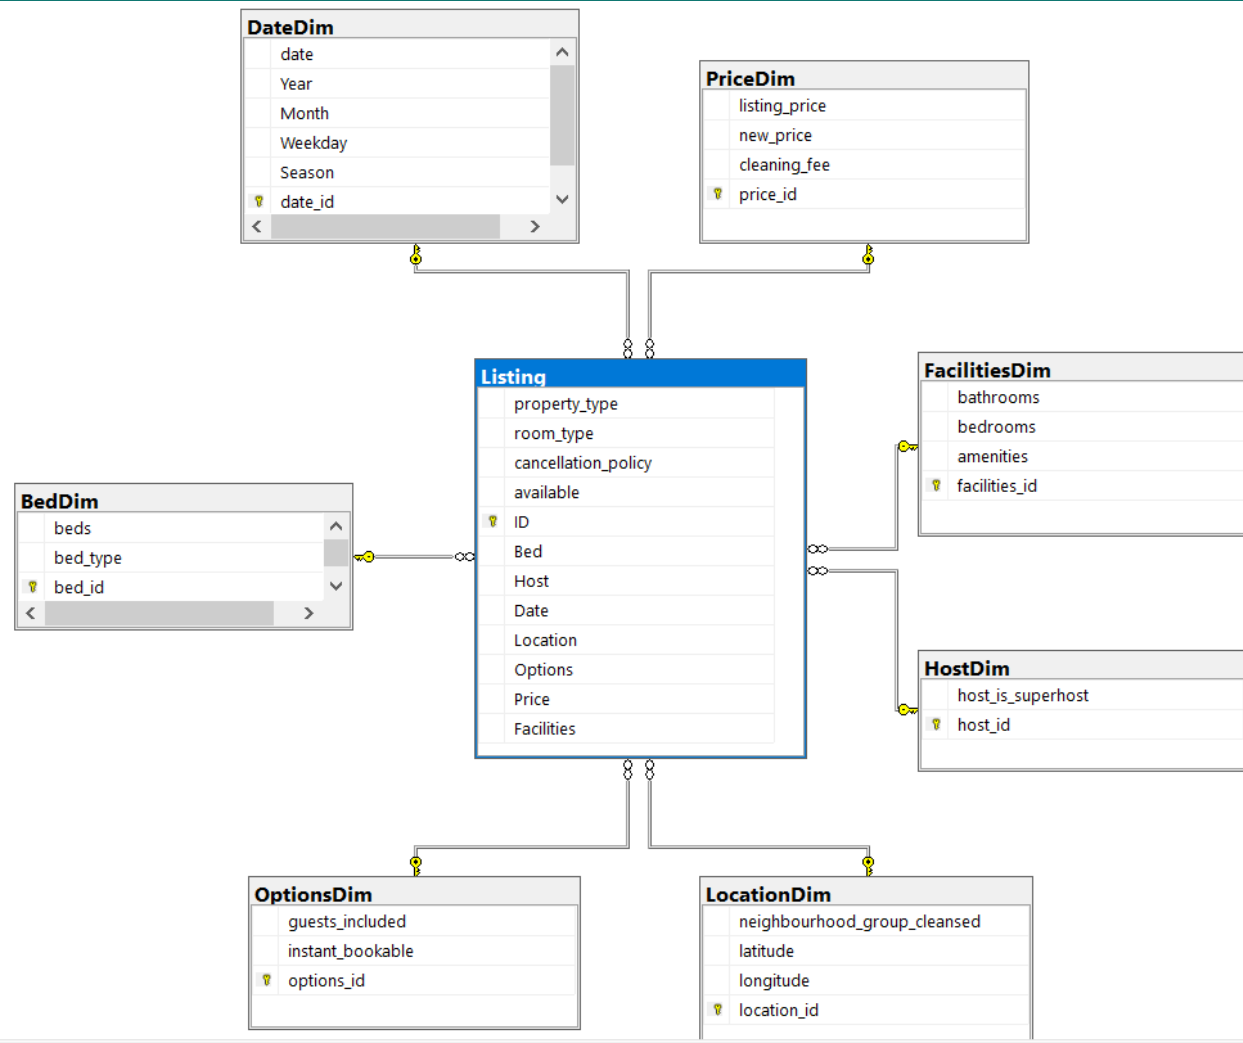
Enforce Foreign Key Cons Yes

► INSERT And UPDATE Spec

Add Delete Close




Star Scheme



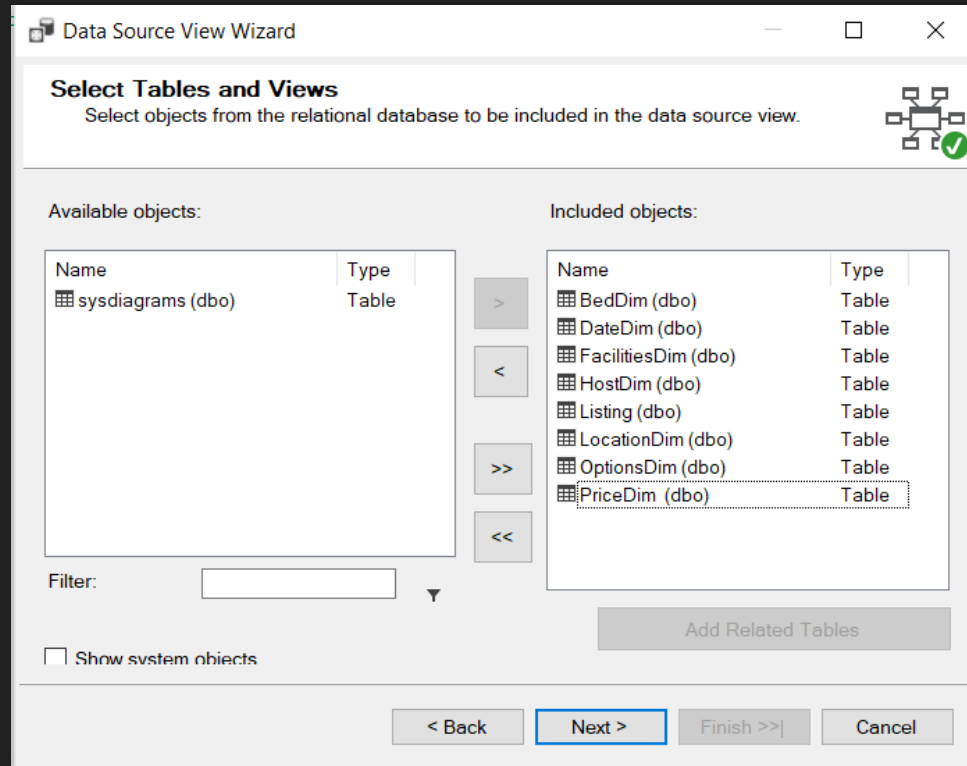
Fact Table – Metrics

Metrics

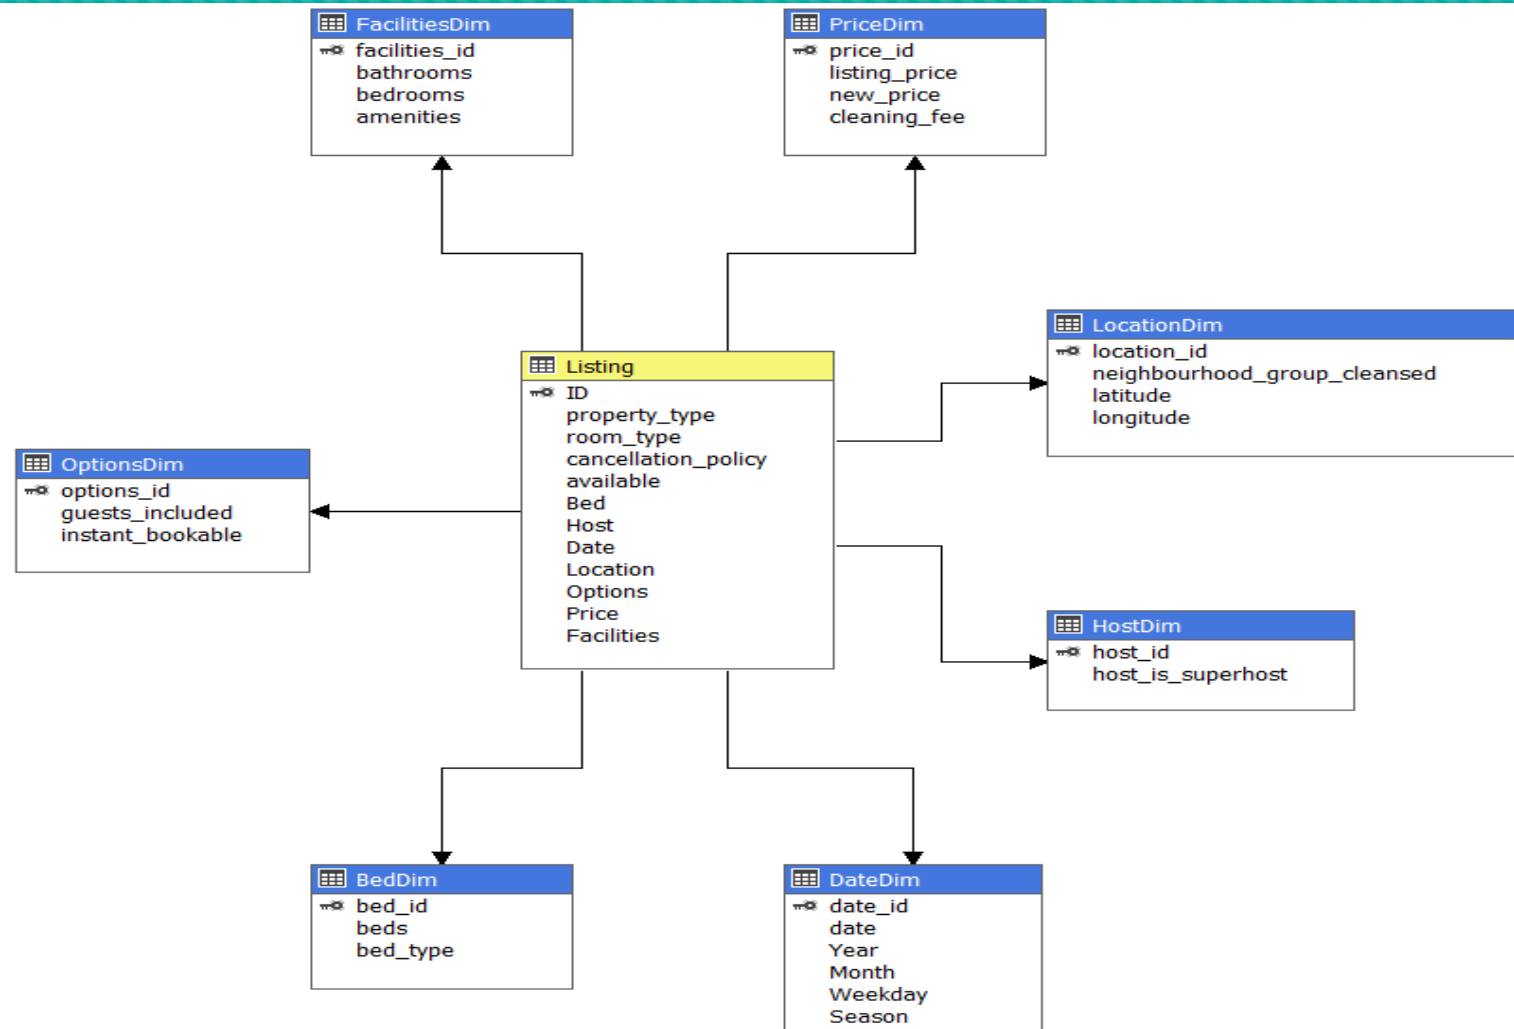


Listing	
	property_type
	room_type
	cancellation_policy
	available
📌	ID
	Bed
	Host
	Date
	Location
	Options
	Price
	Facilities

Data Cube



Data Cube



Calculated Measures


- 2699 Total Listings
- 17 Seattle Neighborhoods referred
- 32% of the Cancellation Policies is Strict
- The Average (mean) Regular Price of the Listings is 137\$
- The Average (mean) Price of the Listing the day the day they were searched for is 150\$
- The Average Price Addition for the Cleaning Fee is 62\$
- More than 2/3 of the Searches for Listings were conducted in Winter

Data Visualization




Power BI Desktop


Which is the most popular bed type offered?



Which is the area with the most listings offered?




Which is the cancellation policy for the majority of the listings?




How does the owner of each listing determine on average their price according to season?

Average of price by Season

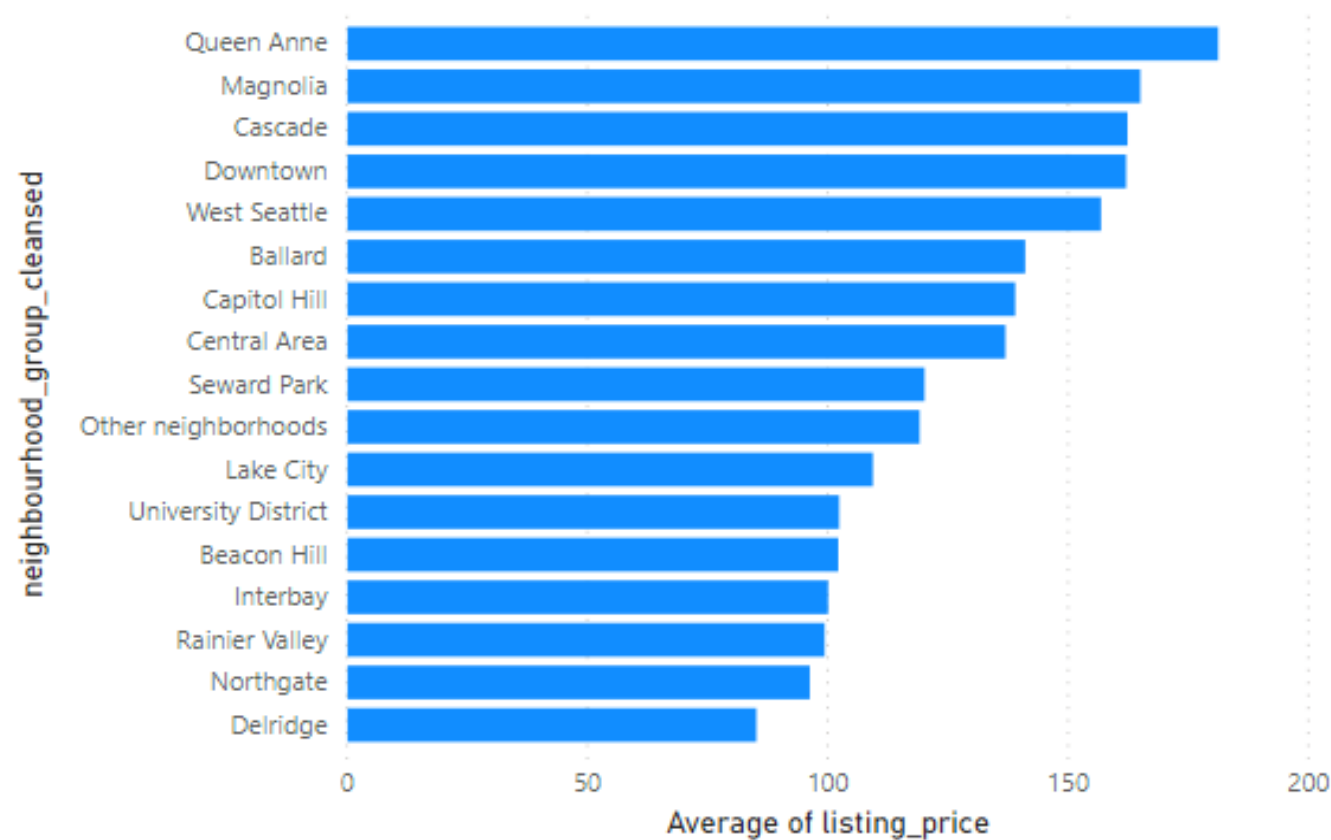


Average of price by Month




How does the owner of each listing determine on average its price according to region?

Average of listing_price by neighbourhood_group_cleansed




Do superhosts exploit their title to charge more?

Average of listing_price by host_is_superhost




How does the host of each listing determine on average its price according to the type of listing he offers?

Average of listing_price by property_type




How different does the average listing price set by the host differ from the average price when adjusted on the basis of season and demand?

Average of listing_price and Average of new_price



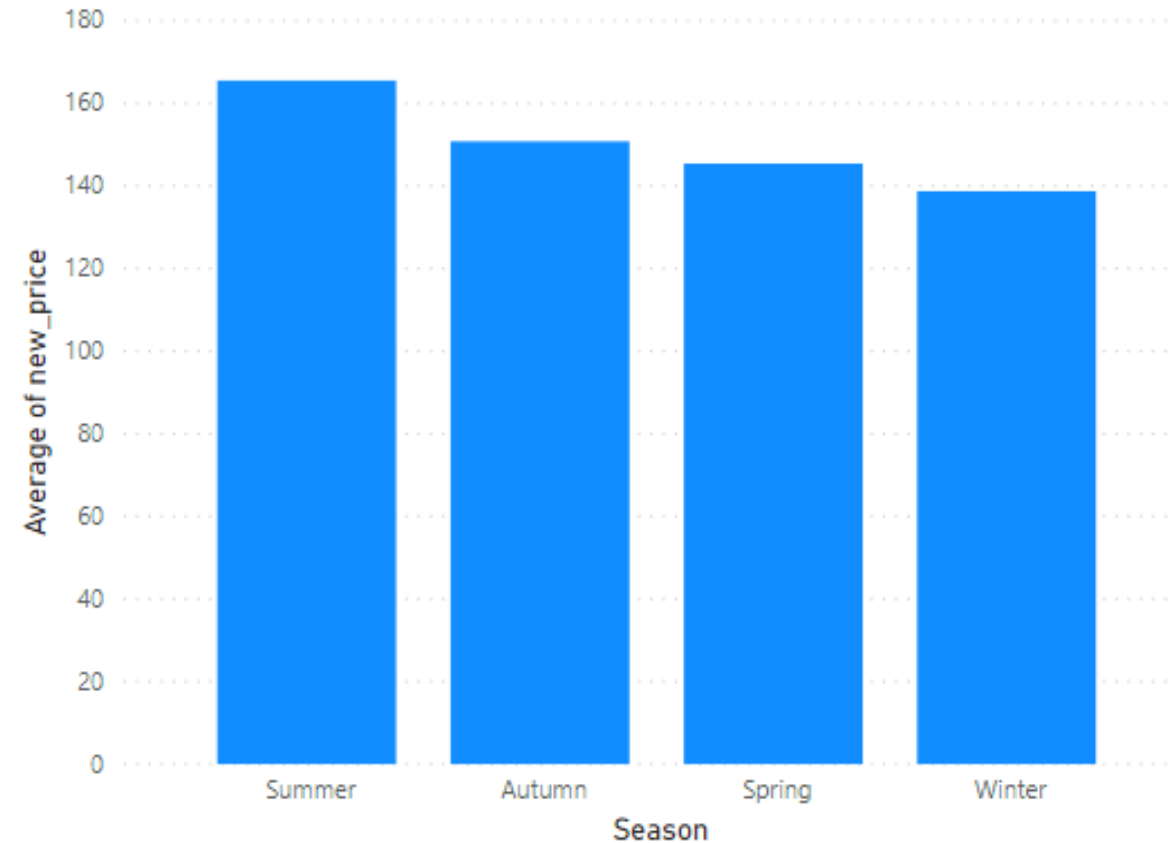
How does each host charge their listing according to the kind of room they offer?

Average of listing_price by room_type




How does each host adapt its listing according to the demand season?

Average of new_price by Season




What is the variance of reviews per month?

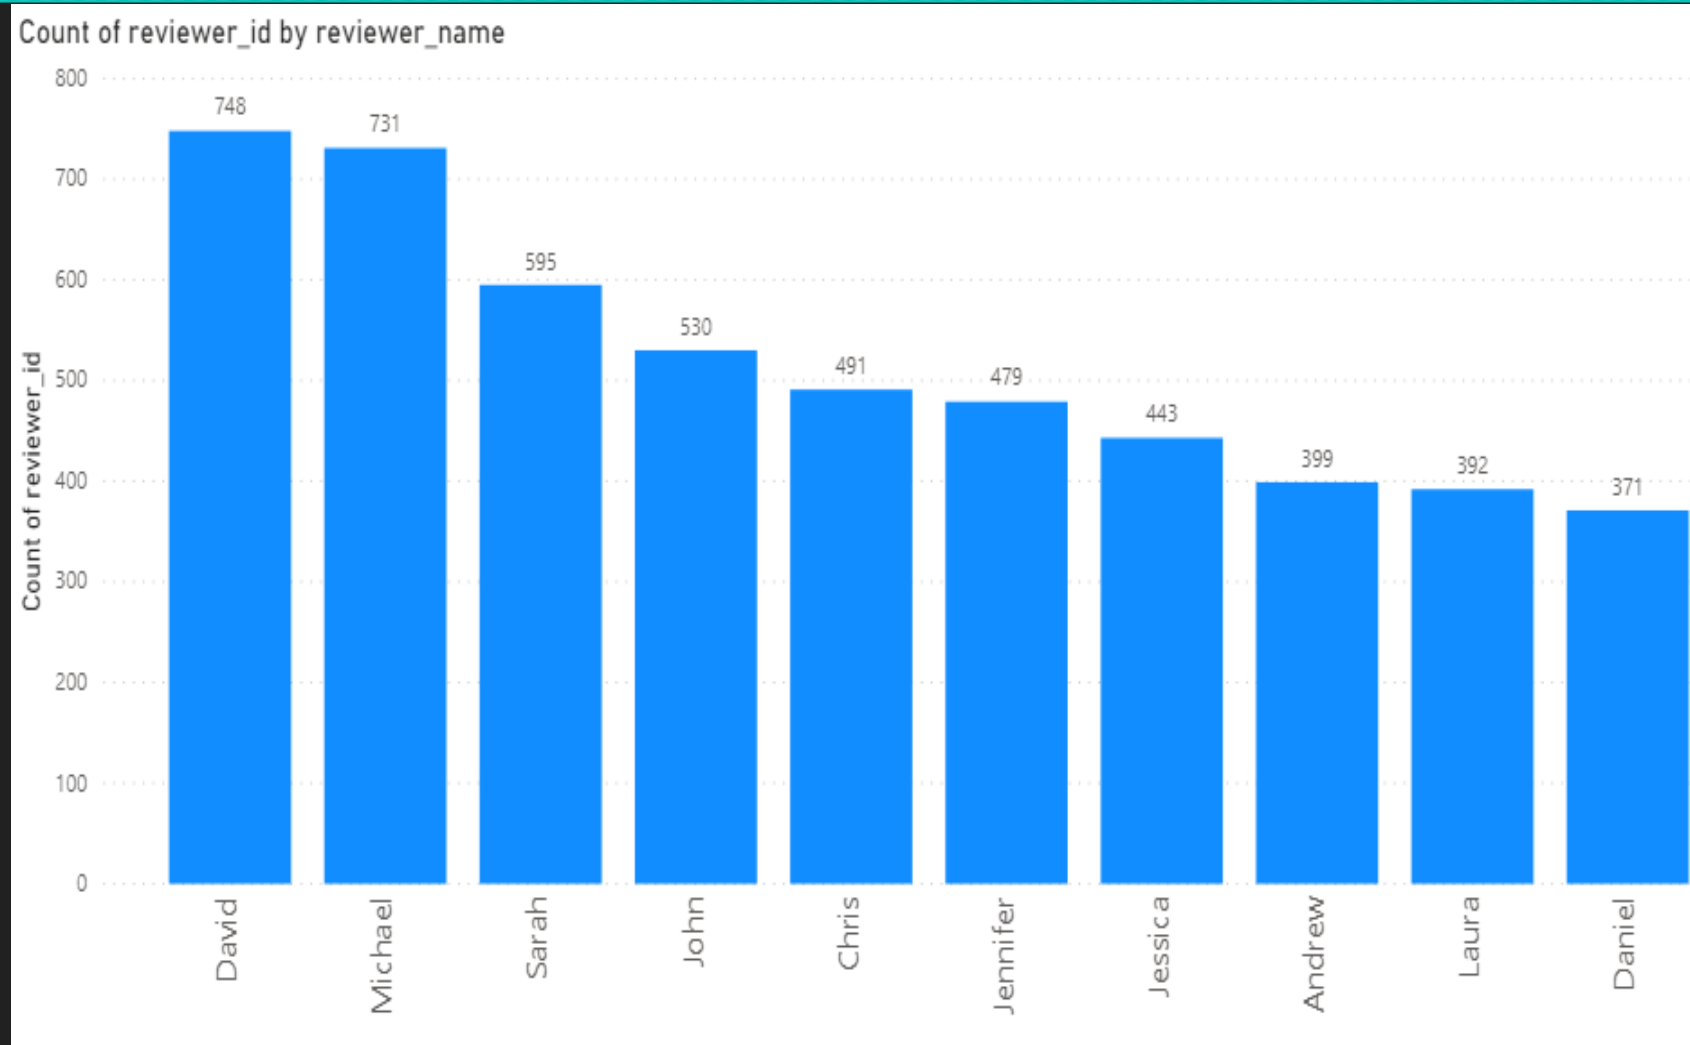
Count of listing_id by Month



Count of listing_id by Year



What's the most common name among tourists?




Classification

Based on the sentiment of the (English reviews)



TOOLS USED: NLTK LIBRARY
FROM PYTHON



NUMBER OF CLASSES:
THREE

Steps before classification

- Remove all the reviews that are in any other language other besides English

```
from langdetect import detect

def detect_lang(sente):
    sente=str(sente)
    try:
        return detect(sente)
    except:
        return "None"

for index,row in reviewsDF.iterrows():
    lang=detect_lang(row['comments'])
    reviewsDF.at[index,'language'] = lang
```

Run the classification

- Run the actual classification with the NLTK library

```
from nltk.sentiment.vader import SentimentIntensityAnalyzer
sid = SentimentIntensityAnalyzer()

reviews_new['polarity_value'] = "Default"
reviews_new['neg'] = 0.0
reviews_new['pos'] = 0.0
reviews_new['neu'] = 0.0
reviews_new['compound'] = 0.0
for index, row in reviews_new.iterrows():
    ss = sid.polarity_scores(row['comments'])
    reviews_new.at[index, 'polarity_value'] = ss
    reviews_new.at[index, 'neg'] = ss['neg']
    reviews_new.at[index, 'pos'] = ss['pos']
    reviews_new.at[index, 'neu'] = ss['neu']
    reviews_new.at[index, 'compound'] = ss['compound']
reviews_new.head()
```

Extract the polarity scores

- Create a dataframe with every information about the reviews and append the columns with the polarity score for each row

	listing_id	id	date	reviewer_id	reviewer_name	comments	polarity_value	neg	pos	neu	com
0	7202016	38917982	2015-07-19	28943674	Bianca	Cute and cozy place. Perfect location to every...	{'neg': 0.0, 'neu': 0.462, 'pos': 0.538, 'comp...	0.000	0.538	0.462	0.79
1	7202016	39087409	2015-07-20	32440555	Frank	Kelly has a great room in a very central locat...	{'neg': 0.0, 'neu': 0.609, 'pos': 0.391, 'comp...	0.000	0.391	0.609	0.98
2	7202016	39820030	2015-07-26	37722850	Ian	Very spacious apartment, and in a great neighb...	{'neg': 0.043, 'neu': 0.772, 'pos': 0.185, 'co...	0.043	0.185	0.772	0.83
3	7202016	40813543	2015-08-02	33671805	George	Close to Seattle Center and all it has to offe...	{'neg': 0.035, 'neu': 0.765, 'pos': 0.2, 'comp...	0.035	0.200	0.765	0.83
4	7202016	41986501	2015-08-10	34959538	Ming	Kelly was a great host and very accommodating ...	{'neg': 0.0, 'neu': 0.655, 'pos': 0.345, 'comp...	0.000	0.345	0.655	0.93

Visualize the classes

- Create 3 different arrays to visualize the polarity scores

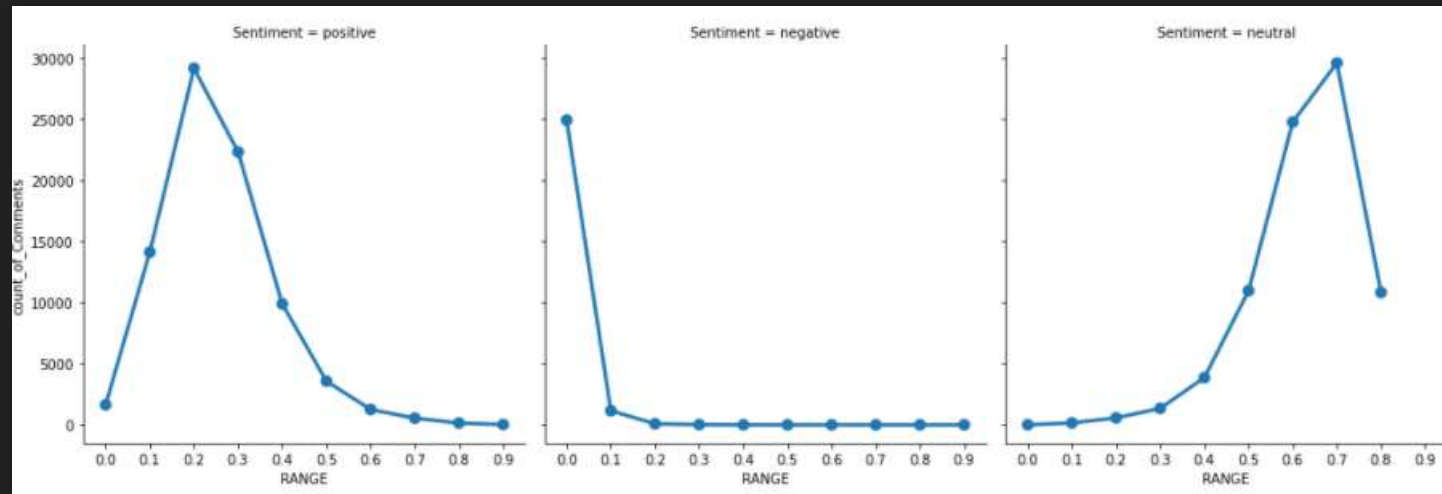
	count_of_Comments	RANGE	Sentiment
0	1640	0.0	positive
1	14126	0.1	positive
2	29151	0.2	positive
3	22306	0.3	positive
4	9881	0.4	positive

	count_of_Comments	RANGE	Sentiment
0	24916	0.0	negative
1	1156	0.1	negative
2	86	0.2	negative
3	21	0.3	negative
4	3	0.4	negative

	count_of_Comments	RANGE	Sentiment
0	8	0.0	neutral
1	157	0.1	neutral
2	569	0.2	neutral
3	1347	0.3	neutral
4	3864	0.4	neutral

Understand the distributions

- Visualize in a better way the polarity scores



Clustering

Based on text descriptions:

id

name

space

description

neighborhood_overview

neighbourhood_cleansed

Clustering



ALGORITHM USED:
K-MEANS



NUMBER OF CLUSTERS:
6

Steps before clustering

Combine column's text into one column

```
# let's combine the name, space, description, and neighborhood_overview into a new column
Listings['combined_description'] = df.apply(lambda x: '{} {} {} {}'.format(x['name'], x['space'],
                                                                           x['description'], x['neighborhood_overview']), axis=1)

print(Listings.loc[0, 'combined_description'])
```

Stylish Queen Anne Apartment Make your self at home in this charming one-bedroom apartment, centrally-located on the west side of Queen Anne hill. This elegantly-decorated, completely private apartment (bottom unit of a duplex) has an open floor plan, bamboo floors, a fully equipped kitchen, a TV, DVD player, basic cable, and a very cozy bedroom with a queen-size bed. The unit sleeps up to four (two in the bedroom and two on the very comfortable fold out couch, linens included) and includes free WiFi and laundry. The apartment opens onto a private deck, complete with it's own BBQ, overlooking a garden and a forest of black bamboo. The Apartment is perfectly-located just one block from the bus lines where you can catch a bus and be downtown Seattle in fifteen minutes or historic Ballard in ten or a quick five-minute walk will bring you to Whole Foods and Peet's Coffee or take a fifteen minute walk to the top of Queen Anne Hill where you will find a variety of eclectic shops, bars, and restaurants. There is no nan

Steps before clustering and Clustering

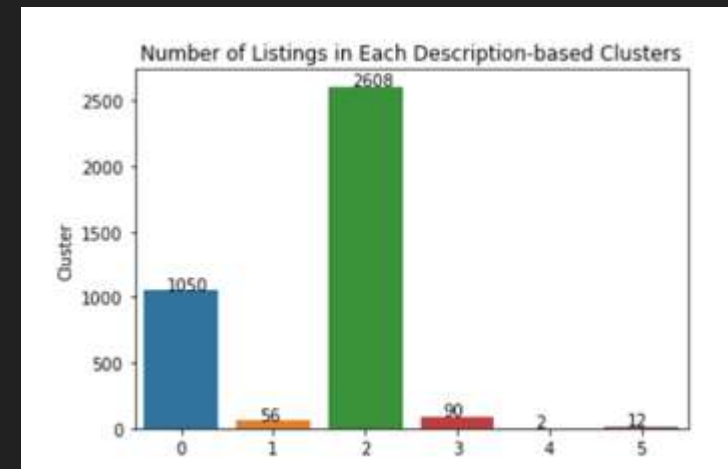
- Set weight for words

```
# Transform combined_description into tfidf format
tfidf = TfidfVectorizer(ngram_range=(1,2), stop_words='english', tokenizer=LemmaTokenizer())
tfidf.fit(df['combined_description'])
DescTfidf = tfidf.transform(df['combined_description'])
```

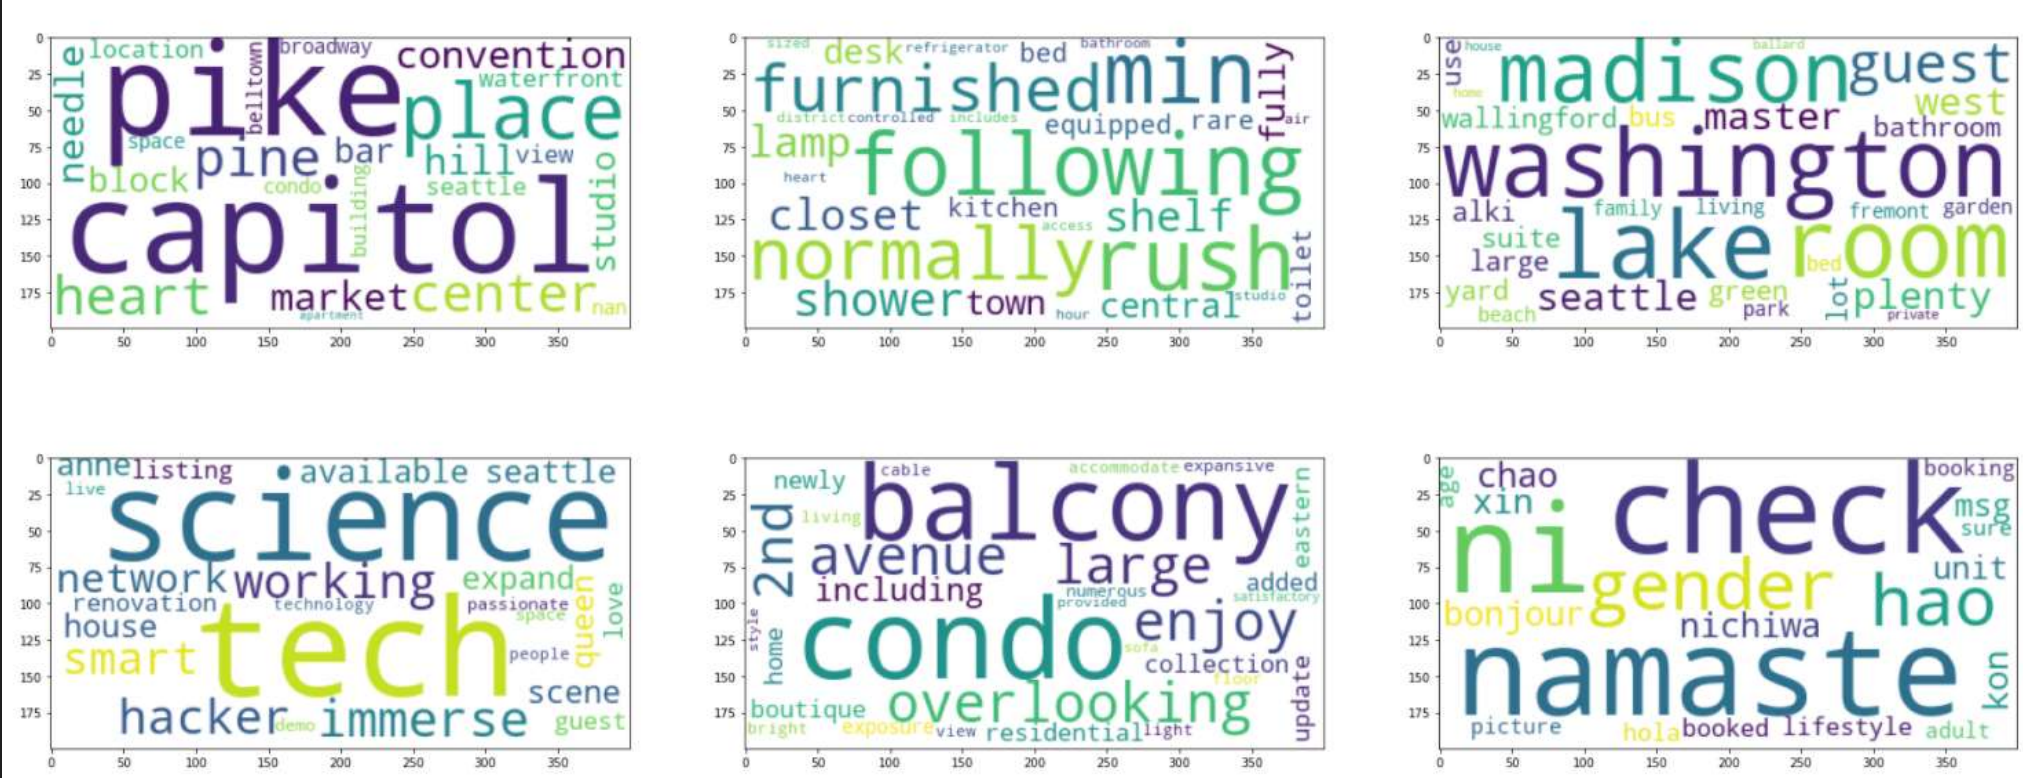
- Keep top 30 words

```
# Top 30 words that describe each cluster
# Pipeline to identify top 30 words that are "best predictor" of a cluster
pipeline = Pipeline([('tfidf', TfidfVectorizer(ngram_range=(1,2), stop_words='english', tokenizer=LemmaTokenizer())),
                    ('clf', SGDClassifier(loss='hinge', penalty='l2',
                                         alpha=1e-3, max_iter=5, random_state=42))),
                    1)
```

- Number of Listings per Cluster



Clusters Overview with wordcloud



Clusters

Cluster 1: «Next to Capitol»



Cluster 2: «Furnished»



Clusters

Cluster 3: «Next to a Lake»



Cluster 4: «Tech House»,



Clusters

Cluster 5: «With a balcony»



Cluster 6: «Chinese»



Prediction

- Predict prices based on

```
df = listings[["host_response_rate", "host_acceptance_rate", "host_is_superhost",  
              "host_listings_count", "zipcode", "property_type", "room_type", "accommodate  
s", "bathrooms", "bedrooms",  
              "beds", "price", "number_of_reviews", "review_scores_rating", "cancellation  
_policy",  
              "reviews_per_month"]]
```

Steps before predicting

Create dummies

```
# select non-numeric variables and create dummies
non_num_vars = df2.select_dtypes(include=['object']).columns
df2[non_num_vars].head()
```

```
# split into test and training data
np.random.seed(1)
indices = np.random.permutation(len(df3))
train_size = int(round(0.8*len(df3)))
test_size = len(df3)-train_size
```

```
y = df3['price']
x = df3.drop('price', axis =1)
```

```
x.train = x.iloc[indices[0:train_size]]
y.train = y.iloc[indices[0:train_size]]
x.test = x.iloc[indices[train_size+1:]]
y.test = y.iloc[indices[train_size+1:]]
```

```
x2 = x.train.as_matrix()
y2 = y.train.as_matrix()
```

Random Forest

```
from sklearn.ensemble import RandomForestRegressor
forest = RandomForestRegressor(n_estimators=500,
                              criterion='mse',
                              random_state=3,
                              n_jobs=-1)


forest.fit(X_train, y_train)
y_train_pred = forest.predict(X_train)
y_test_pred = forest.predict(X_test)

print('MSE train: %.3f, test: %.3f' % (
    mean_squared_error(y_train, y_train_pred),
    mean_squared_error(y_test, y_test_pred)))
print('R^2 train: %.3f, test: %.3f' % (
    r2_score(y_train, y_train_pred),
    r2_score(y_test, y_test_pred)))
```

```
MSE train: 360.229, test: 2275.514
R^2 train: 0.945, test: 0.670
```

Plot results

```
plt.scatter(y_train_pred, y_train_pred - y_train,  
            c='blue', marker='o', label='Training data')  
plt.scatter(y_test_pred, y_test_pred - y_test,  
            c='lightgreen', marker='s', label='Test data')  
plt.xlabel('Predicted values')  
plt.ylabel('Residuals')  
plt.legend(loc='upper left')  
plt.show()
```



Any Questions?

Thank you